

Causal-Factored Cross-City Model Transfer for Bus Holding Control with Open Transit Data

Author names withheld for review

Abstract

Hybrid offline-to-online reinforcement learning for transit control is attractive because operators often have rich historical data but only imperfect simulators. However, a bus holding policy tuned in one city may fail in another city unless the target simulator is manually calibrated, and such calibration is expensive. We study this cross-city transfer problem for bus holding using four open transit city bundles: Singapore, Austin/CapMetro, Halifax Transit, and MBTA. We propose Causal-Factored Cross-City Model Transfer (CFCMT), a mechanism-wise residual world-model adapter that decomposes dense simulator corrections into headway, demand/dwell, speed, and reward mechanisms, and optionally reweights source cities by static speed, headway, route, and demand similarity. Against a dense $H2O^+$ -style residual baseline, CFCMT is evaluated under a strict leave-one-city-out protocol in which each of the four cities is the target exactly once. On static-derived cross-city dynamics validation, unweighted CFCMT beats $H2O^+$ on all four targets with mean total-MSE ratio 0.759, while similarity-weighted CFCMT also wins on all four targets and lowers the mean ratio to 0.578. In one-step policy validation, CFCMT improves reward/regret on three of four targets and lowers bunching rate on all four targets. Calibration-budget, per-mechanism, ablation, route-bootstrap, source-subset, and generator-bias checks support the same main trend while showing that the advantage depends on multi-source diversity and should not be overstated as superiority on every small city subset. On two public Austin raw APC/AVL slices, source-only CFCMT is more stable than dense residual transfer but still worse than a state-informed passive no-correction baseline; with few-shot target residual calibration, a residual gate beats that passive baseline at 1% target line-group budget in both slices. The resulting contribution is a reproducible, reviewer-facing protocol and an empirically supported boundary condition: mechanism-factored transfer is useful for low-calibration cross-city residual adaptation, but counterfactual control or field validation remains necessary before operational claims.

1 Introduction

Bus holding control aims to reduce headway instability by delaying selected vehicles at stops. The control problem is local in execution but system-wide in consequence: a hold action changes forward and backward headways, passenger waiting, dwell, and downstream bunching. Modern reinforcement learning (RL) methods can learn holding policies in simulation, but practical deployment is constrained by two persistent gaps. First, high-fidelity simulation and manual calibration are expensive to reproduce in each city. Second, public data availability differs sharply across agencies: one city may expose speed bands and passenger OD data, while another may only provide GTFS schedules and coarse ridership summaries.

Hybrid offline-to-online methods such as $H2O^+$ learn from both logged data and simulator interaction by correcting or reweighting dynamics gaps (Niu et al., 2022, 2025). This is a natural starting point for transit systems, where a historical buffer and a lower-fidelity simulator are often available. The difficulty is cross-city transfer. A dense monolithic residual model can fit simulator

error in the source cities, but its parameters mix mechanisms that may shift differently across cities: schedule headways, route speeds, stop demand, and dwell behavior do not all change together. If cross-city transfer depends on localized changes in mechanisms, a factored representation should be more stable than a dense one (Bengio et al., 2020; Feng et al., 2022).

This paper studies that hypothesis in a deliberately conservative setting. We do not claim that the current data are vehicle-level real-world trajectories. Instead, we construct full-route open transit city bundles and evaluate models on static-data-derived offline dynamics targets. The protocol is meant to answer a narrower question: when only open static data and an uncalibrated simulator-like transition generator are available, does a mechanism-factored residual adapter transfer better across cities than a dense $H2O^+$ -style residual correction?

We introduce Causal-Factored Cross-City Model Transfer (CFCMT). CFCMT fits residual corrections at the level of output mechanisms rather than one dense residual over all outputs. For each target quantity, it uses a mechanism-specific sparse parent set: headway and gap outputs depend primarily on neighboring headways, target headway, speed, time, and action; demand and dwell outputs depend on waiting passengers, speed, target headway, time, and action; speed outputs use speed, time, and route-position features; reward uses policy-relevant headway and waiting quantities. We also add a city-similarity source weighting rule that uses only static descriptors available before target calibration.

The paper contributes:

1. A cross-city open-transit validation protocol using four full-route city bundles and a strict leave-one-city-out evaluation unit.
2. A mechanism-wise residual transfer model and a static descriptor based source weighting rule.
3. A set of reviewer-facing diagnostics: source-weight sensitivity, source-subset robustness, generator-bias/noise robustness, and traffic-domain policy metrics.
4. An explicit boundary on the claim: the evidence supports static-derived cross-city transfer and not yet operational deployment without vehicle-level validation.

2 Related Work

Bus holding control. Classical holding policies analyze headway stability and bunching through control rules and queueing approximations (Daganzo, 2009; Xuan et al., 2011). Recent work has applied deep RL and multi-agent RL to bus holding under stochastic demand and travel times (Alesiani and Gkiotsalitis, 2018; Wang and Sun, 2020; Liu et al., 2023; Wang and Sun, 2023). These studies typically focus on one calibrated network or one simulation benchmark. Our focus is different: we evaluate whether a learned dynamics correction transfers across city bundles with heterogeneous open data sources.

Hybrid offline-to-online RL and dynamics gaps. $H2O^+$ and $H2O^+$ -style methods address the mismatch between offline data and online simulator interaction by estimating dynamics-related weights or residuals (Niu et al., 2022, 2025). Related sim-to-real and offline-to-online methods use domain classifiers, calibrated Q-learning, residual dynamics, or world-model adaptation to reduce transfer error (Eysenbach et al., 2021; Nakamoto et al., 2023; Lanier et al., 2025). In this paper, $H2O^+$ is treated as the dense residual-correction baseline: it can learn flexible source-city corrections, but the correction is not explicitly decomposed by transit mechanism.

Causal and factored transfer. The independent causal mechanisms principle suggests that distribution shifts may be localized when represented in appropriate mechanisms (Bengio et al., 2020). Factored adaptation methods use this idea to improve non-stationary RL and transfer by separating latent change factors or identifying the variables that need target-domain adaptation (Huang et al., 2022; Feng et al., 2022). CFCMT applies this principle to bus holding dynamics: rather than learning one dense correction, it learns mechanism-wise residuals aligned with interpretable transit outputs. The approach is intentionally lightweight in this paper: the reported cross-city experiments use sufficient-statistics ridge models for reproducibility, while the repository also contains neural causal-factored residual modules for later end-to-end training.

Traffic simulation and open transit data. SUMO provides a standard microscopic traffic simulation platform (Lopez et al., 2018). For cross-city public transit studies, however, fully calibrated SUMO networks and vehicle-level AVL/APC data are rarely available across many agencies. We therefore use open static data to create full-route city bundles. This makes the experiments reproducible but also limits their interpretation; the static-derived targets are not a substitute for live operational trajectory validation.

3 Problem Setting

Let $s_t \in \mathcal{S}$ denote a bus-system state at a decision event and $a_t \in \mathcal{A}$ denote a holding action in seconds. The state contains line and station identifiers, time of day, direction, forward and backward headways, waiting passengers, target headway, base dwell, co-line headways, gap, and segment speed. The transition target is

$$\mathbf{y}_t = (h_{t+1}^+, h_{t+1}^-, q_{t+1}, d_{t+1}, g_{t+1}, \bar{h}_{t+1}^+, \bar{h}_{t+1}^-, v_{t+1}, r_{t+1}),$$

covering next headways, waiting passengers, dwell, gap, co-line headways, speed, and reward.

Each city c provides a bundle $\mathcal{D}_c$ of route, timetable, speed, and demand tables. From these tables we construct paired simulator and static-derived target transitions:

$$\mathbf{y}^{\text{sim}} = f_0(s, a), \quad \mathbf{y}^{\text{target}} = f_c(s, a),$$

and train residual models on

$$\Delta_c(s, a) = \mathbf{y}^{\text{target}} - \mathbf{y}^{\text{sim}}.$$

For cross-city validation, source cities $\mathcal{C}_{\text{src}}$ are used for fitting and a held-out target city c^* is used only for evaluation.

The main evaluation unit is a city target, not a route. For four cities, the strict leave-one-city-out protocol has exactly four target evaluations:

$$\{(\mathcal{C} \setminus \{c\}) \rightarrow c : c \in \mathcal{C}\}.$$

This avoids counting repeated target cities as independent evidence.

Claim scope. The target transitions are static-data-derived pseudo-observations. They combine agency data, schedule-derived speeds or headways, passenger-demand summaries, and deterministic transition rules. They are useful for reproducible cross-city validation, but they are not vehicle-level AVL/APC ground truth. Accordingly, the paper evaluates transfer robustness under open static-data-derived dynamics, not operational field performance.

4 Method

4.1 Dense H2O⁺-style residual baseline

The H2O⁺-style baseline fits one dense residual model over all transition outputs:

$$\hat{\Delta}_{\text{dense}}(s, a) = \phi_{\text{dense}}(s, a)^\top B,$$

where ϕ_{dense} includes normalized state/action features, squares, action interactions, and selected pairwise products. At evaluation time the corrected prediction is

$$\hat{y}_{\text{H2O}^+}(s, a) = f_0(s, a) + \hat{\Delta}_{\text{dense}}(s, a).$$

The baseline is intentionally strong: it can use many cross terms and is fit on all source-city transitions.

4.2 Causal-factored residual transfer

CFCMTdecomposes the residual by output mechanism:

$$\hat{\Delta}_j(s, a) = \phi_j(s, a)^\top \beta_j, \quad j \in \{1, \dots, 9\}.$$

The feature map ϕ_j is chosen according to the output mechanism. Headway and gap residuals use current headways, target headway, speed, action, and time. Waiting and dwell residuals use waiting passengers, target headway, speed, station fraction, action, and time. Speed residuals use current speed, event time, and station position. Reward residuals use policy-relevant headway deviations, waiting passengers, gap, speed, action, and time. The corrected prediction is

$$\hat{y}_{\text{CFCMT}}(s, a) = f_0(s, a) + \left(\hat{\Delta}_1(s, a), \dots, \hat{\Delta}_9(s, a) \right).$$

All reported models use ridge-regularized sufficient statistics, so the experiment can be reproduced without stochastic neural training.

4.3 Source-city similarity weighting

For each city we compute a static descriptor z_c from speed quantiles, headway quantiles, route-distance quantiles, peak-demand hour, and the 24-hour demand-share vector. For a target city t and source city c , the distance is computed after feature-wise standardization over the source-plus-target set:

$$d(c, t) = \frac{1}{\sqrt{m}} \left\| \frac{z_c - \mu}{\sigma} - \frac{z_t - \mu}{\sigma} \right\|_2.$$

The source weight is

$$w_c = \frac{\exp(-d(c, t)/T) + \lambda \max_{c'} \exp(-d(c', t)/T)}{\sum_{k \in \mathcal{C}_{\text{src}}} [\exp(-d(k, t)/T) + \lambda \max_{c'} \exp(-d(c', t)/T)]}.$$

The default values are $T = 1.0$ and $\lambda = 0.05$. The source sufficient statistics are rescaled so that these weights control the effective number of source transitions rather than simply reflecting raw city size.

city	em_jer	dominal_jer	traffic_jer	line	stop	segregate	transferable	dominal_jer2	dominal_jer3	dominal_jer4	peak_dominal_jer	peak_dominal_jer2	traffic_crowd	
Singapore	singapore_jer1	sharpened stop OD/PM monthly	sharpened speed bands	70	29673	2090	50515	124188	17.0	3.58	600.0	7	LTA Roadmap PM/OD/20th WEEKDAY assembly totally divided to typical weekday	LTA Trafficflow/pollution/segregate assestment matched to stop segregate
Atlanta / Capetown	atlanta_jer2	schedule stop	schedule speed	24	6864	604	28329	33072	15.783	6.19	1200.0	14	GTPS schedule busway segment OD, on public use stop ramped bus	GTPS schedule-divided segment speed
London	london_jer1	stop	stop	1728	1818	100	100	11800	11800	11800	11800	15	Schedule busway bus-lane/stop on GTPS road network	GTPS schedule-divided segment speed
MRTA	mrtajer1	stop board/aligned operation	schedule divided	2035	2035	2035	6917	90930	7.6405	4.01	100.0	7	MRTA PM 2021 speed of boarding/aligned operation across GTPS route-direction pattern	GTPS schedule-divided segment speed

Table 1: Open transit city bundles. Evidence levels distinguish observed passenger/traffic data from schedule-derived proxies.

4.4 One-step policy validation

For policy-level validation, each method evaluates a discrete hold set $\{0, 15, 30, 45, 60\}$ seconds under a set of headway-stress states. The policy selects the action with the highest predicted reward under its corrected transition model. The selected action is then scored against the static-derived target transition. We report reward, regret to an oracle action, mean absolute headway error, bunching rate, large-gap rate, and mean hold seconds. This is a policy-level one-step lookahead benchmark, not a live SUMO rollout.

5 Experiments

5.1 Data

We evaluate on four full-route city bundles. Singapore uses LTA DataMall passenger-volume and OD bus data with traffic speed bands. Austin/CapMetro uses GTFS schedules with demand proxied from scheduled headways. Halifax Transit uses route-level APC summaries apportioned over GTFS route-direction patterns. MBTA uses Fall 2025 stop-level boardings and alightings apportioned over GTFS route-direction patterns. Table 1 summarizes the bundles and their evidence levels.

5.2 Validation protocols

Primary dynamics validation. The main dynamics experiment is strict leave-one-city-out. Each city is held out once as the target, and all remaining cities are used as sources. The metric is total MSE averaged over the nine transition outputs. Ratios below 1 indicate improvement over the dense H2O⁺-style residual baseline.

Policy validation. The one-step policy experiment uses the same source-target splits and evaluates action selection under headway-stress states. We report reward/regret and traffic-domain metrics. Bunching rate is the fraction of decisions in which either adjacent headway is below 0.5 times the target headway. Large-gap rate is the fraction in which either adjacent headway exceeds 1.5 times the target headway.

Reviewer-facing robustness checks. We include diagnostics targeted at common review objections. First, a target-construction audit separates observed, apportioned, schedule-proxy, and deterministic target components. Second, a target-route calibration-budget sweep adds 0%, 1%, 5%, 10%, 25%, or 100% of held-out target-city routes to the source sufficient statistics and evaluates on the remaining target routes; the 100% point is an in-sample oracle reference. Third, per-mechanism errors decompose total MSE into headway/gap, demand/dwell/reward, and speed groups. Fourth, mechanism ablations compare the proposed parent sets against wrong mechanism groupings, action/time-only features, and dense matched-capacity variants. Fifth, source-weight sensitivity varies $T \in \{0.5, 1.0, 2.0\}$ and $\lambda \in \{0, 0.05, 0.1\}$, while source-subset and pairwise-source checks remove major source families. Sixth, generator robustness perturbs output groups or adds measurement-noise terms to the evaluation sufficient statistics.

experiment	paper_role	unit	scope_note
strict_leave_one_city_out	primary cross-city dynamics evidence	city target	four unique targets; no duplicated target counted as independent
generator_robustness	bias/noise defense	scenario x city target	static-derived target perturbation; not real AVL/APC trajectory ground truth
source_subset_robustness	data-source confound defense	city subset	reports excluding Singapore and excluding Austin proxy-demand subsets
cross_city_policy_validation	primary one-step policy evidence	city target	one-step lookahead, not live SUMO rollout
sampled_rollout	auxiliary sanity check only	BusSimEnv episode	50 policy episodes across runnable line envs; do not use as primary claim

Table 2: Experimental scope and paper role. The primary evidence is strict city-level leave-one-out dynamics and one-step policy validation.

target_city	source_cities	target_transitions	h2oplus_total_mse	cfcmt_total_mse	weighted_cfcmt_total_mse	cfcmt_vs_h2oplus_ratio	weighted_cfcmt_vs_h2oplus_ratio	weighted_cfcmt_vs_unweighted_ratio
Singapore	austin_capmetro_all, halifax_transit_all, mbta_all	1243488	2254.8548	2246.6693	2251.7081	0.9964	0.9986	1.0022
Austin / CapMetro	singapore_jta_all, halifax_transit_all, mbta_all	310272	91.9724	67.8254	36.6531	0.7375	0.3985	0.5404
Halifax Transit	singapore_jta_all, austin_capmetro_all, mbta_all	266448	97.1175	69.5992	54.7536	0.7166	0.5638	0.7867
MBTA	singapore_jta_all, austin_capmetro_all, halifax_transit_all	969360	318.9925	186.4177	112.2698	0.5844	0.352	0.6922

Table 3: Strict leave-one-city-out dynamics results. Ratios below 1 indicate lower total MSE than H2O⁺.

External real APC/AVL validation. To partially address the static-derived target limitation, we add an external Austin/CapMetro validation slice from public raw APC stop-event data. The July–December 2021 and August 2016 files contain observed stop events with boardings, alightings, vehicle load, dwell time, vehicle position, route, trip, stop, and vehicle identifiers. We filter records with valid stop, trip, GPS, and import-quality flags, reconstruct stop-level headways by route-direction-stop event times, and evaluate passive action-0 one-step dynamics against the observed next stop event. The no-correction baseline in this experiment is a state-informed passive transition using the APC-derived current state, not a schedule-only free-running simulator. This is real historical APC/AVL validation, but it is still offline passive validation rather than a field control deployment because counterfactual hold actions are not observed. We also run a few-shot APC adaptation sweep at 0%, 1%, 5%, 10%, and 25% of Austin line groups. This sweep compares source-only transfer, source-plus-target residual refitting, target-only residual fitting, a residual gate that learns how much of the source correction to trust, and a target residual-bias adapter. The gate and bias-adapter variants are lightweight AdaRL/FANsRL-style checks of what changes should be adapted; they are not full reimplementations of those neural RL systems.

Auxiliary executable rollout. Sampled BusSimEnv rollout is reported only as an auxiliary sanity check because only Austin and Halifax have runnable generated line environments. For this rollout, each policy is evaluated on the same line-level random seed, including Python signal-delay randomness, and live observations are mapped back to the static-transition feature schema before model scoring. The policy score uses the BusSimEnv linear headway reward rather than the static one-step reward.

6 Results

6.1 Strict leave-one-city-out dynamics

Figure 1 and table 3 show the primary dynamics result. Unweighted CFCMT beats the dense H2O⁺-style residual baseline on all four held-out cities, with mean total-MSE ratio 0.759. Similarity-weighted CFCMT also wins on all four targets and improves the mean ratio to 0.578. The largest gains occur for Austin, Halifax, and MBTA. Singapore is the hardest target: unweighted CFCMT is only slightly better than H2O⁺, and similarity weighting is slightly worse than unweighted CFCMT while still remaining below the H2O⁺ baseline.




Figure 1: Strict leave-one-city-out dynamics validation. Bars show total-MSE ratio relative to $H2O^+$. The dashed line is parity with $H2O^+$.




Figure 2: Strict leave-one-city-out one-step policy regret ratio. Ratios below 1 favor CFCMT over $H2O^+$.

6.2 Policy-level validation

The strict one-step policy summary is more mixed but still favorable. CFCMT improves reward/regret on three of four targets, improves mean absolute headway error on three of four, and lowers bunching rate on all four targets. The mean regret ratio is 0.876, and the mean bunching-rate ratio is 0.644. Halifax is the main policy exception: CFCMT lowers headway error slightly but has worse reward/regret than $H2O^+$. This is consistent with a limited policy claim: mechanism-factored transfer improves cross-city dynamics strongly and often improves one-step policy quality, but it is not a universal policy dominance result.

6.3 Source weighting is stable under a small grid

The default source weighting uses $T = 1.0$ and $\lambda = 0.05$. Across the 3×3 grid, all nine settings beat $H2O^+$ on all six configured cross-city splits. The mean weighted-CFCMT/ $H2O^+$ total-MSE ratio ranges only from 0.6025 to 0.6149. The best setting in the grid is $T = 0.5, \lambda = 0$, but the

split	target_city	h2o_plus_reward	cfcmt_reward	cfcmt_reward_gain	cfcmt_regret_ratio	cfcmt_headway_error_ratio	h2o_plus_bunching_rate	cfcmt_bunching_rate	cfcmt_bunching_ratio	h2o_plus_large_gap_rate	cfcmt_large_gap_rate	cfcmt_large_gap_ratio	cfcmt_hold_time_delta
leave-one-city-out-singapore	Singapore	-0.3338	-0.3337	0.0001	0.9859	1.0007	0.0024	0.0024	0.9991	0.0016	0.0016	1.0	-0.3597
leave-one-city-out-singapore-halifax	Halifax / CapMetro	-0.3338	-0.332	0.0017	0.975	0.9995	0.0	0.0	0.9991	0.0001	0.0	0.0	0.3999
leave-one-city-out-halifax-halifax	Halifax Transit	-0.3307	-0.3312	-0.0005	1.1113	0.9998	0.0	0.0	0.9355	0.0	0.0	1.0	1.4814
leave-one-city-out-mumbai	MHTA	-0.3429	-0.338	0.0049	0.0201	0.9987	0.0006	0.0004	0.6396	0.0008	0.0002	0.3174	-1.1321

Table 4: Strict leave-one-city-out one-step policy metrics. The table includes reward, regret ratio, headway error ratio, bunching, large-gap rate, and hold-time difference.




Figure 3: Source-weighting sensitivity. Each cell reports mean total-MSE ratio versus $H2O^+$ across configured cross-city splits.

default is not selected by post-hoc optimization and remains within the same narrow band.

6.4 Calibration budget and mechanism-level errors

The target-route calibration-budget sweep tests whether a calibrated dense residual baseline quickly erases the zero-calibration CFCMT advantage. At 0% target calibration, weighted CFCMT has mean ratio 0.578 against the dense source-only baseline. When 1%, 5%, 10%, and 25% of target routes are added to the dense baseline, the uncalibrated weighted-CFCMT/ $H2O^+$ ratio remains below 1 on three of four strict targets, with mean ratios 0.581, 0.586, 0.617, and 0.677, respectively. The 100% target-route point is an in-sample oracle reference and should not be interpreted as deployable zero-calibration evidence.

Per-mechanism errors show that the total-MSE advantage is not driven by one output group only. Weighted CFCMT improves all three mechanism groups on average: headway/gap ratio 0.574, demand/dwell/reward ratio 0.771, and speed ratio 0.668. Headway/gap is the largest contributor to the main gain, which is consistent with the holding-control objective.

6.5 Robustness and boundary conditions

Generator robustness supports the main trend but also identifies failure modes. Weighted CFCMT beats $H2O^+$ on every city target in five of seven perturbation/noise scenarios. Under headway-gap bias and mixed-mechanism bias, the result weakens to three of four targets. Thus the evidence does not say that CFCMT is invariant to any target-generation change; it says that the result is robust to several plausible perturbations.

Mechanism ablations support the mechanism-specific interpretation, with a caveat. Wrong mechanism groupings and action/time-only features are much worse than $H2O^+$, with mean ratios




Figure 4: Target-route calibration-budget sensitivity. Ratios compare weighted CFCMT with no target calibration or the same target budget against a dense H2O⁺-style baseline using that target budget.




Figure 5: Per-mechanism mean MSE ratios under strict leave-one-city-out dynamics validation.

above 7. Dense matched sparse features are close to full CFCMT, with mean ratio 0.772, which means the result should be described as evidence for mechanism-wise sparse transfer rather than definitive causal identification.

The source-subset check is the strongest limitation. With all four cities, weighted CFCMT has mean ratio 0.578. Excluding Austin remains favorable with mean ratio 0.663. Excluding Singapore makes the result nearly tied with H2O⁺, with mean ratio 0.998. The two-city Halifax-MBTA pair is slightly worse than H2O⁺, with ratio 1.042. This indicates that source diversity matters: the method should be framed as multi-source cross-city transfer, not as guaranteed transfer between every pair of cities.

Route-level bootstrap confidence intervals reinforce the same boundary. Unweighted CFCMT has a route-bootstrap ratio confidence interval below 1 for all four targets. Weighted CFCMT has intervals below 1 for Austin, Halifax, and MBTA, but Singapore crosses parity with interval [0.993, 1.002]. The conservative reading is that source weighting improves the mean strongly, while Singapore remains the least secure target.




Figure 6: Generator-bias and measurement-noise robustness under strict leave-one-city-out evaluation.

6.6 External real APC/AVL validation

The public Austin raw APC validation is the strongest non-static check in the current paper. Across two real APC/AVL datasets, the evaluation uses 11.57 million observed stop-event transitions after quality filtering. Weighted CFCMT beats the dense H2O⁺-style residual transfer baseline in both years: the total-MSE ratio is 0.330 on July–December 2021 and 0.323 on August 2016. Unweighted CFCMT also beats the dense baseline, with ratios 0.528 and 0.517. This supports the mechanism-factored transfer claim on real historical APC/AVL observations.

The result should still be read conservatively. The APC-derived no-correction passive baseline has lower total MSE than both residual-transfer models on this real APC slice. This does not mean that a free-running uncalibrated simulator is operationally better; it means that the baseline is given the observed current APC state and that source-only residual corrections can overshoot under real operational noise. The overshoot is concentrated in headway/gap outputs, which dominate the unnormalized total MSE, while CFCMT is better than the passive baseline on several demand, dwell, and speed outputs. Thus the real APC result strengthens the comparison between CFCMT and dense residual transfer, but it also shows that a safe residual gate or small target-city calibration budget is needed before claiming improvement over a state-informed passive baseline.

The few-shot APC sweep confirms this diagnosis. With only 1% of Austin line groups used for target residual calibration, the weighted-CFCMT residual gate improves over the passive baseline in both APC datasets, with ratios 0.970 and 0.957. At 25% target line-group budget, target-only CFCMT reaches the best ratios, 0.893 on July–December 2021 and 0.879 on August 2016. Thus the real APC gap is not evidence against the mechanism-factored representation; it is evidence that source-only residual transfer needs target-domain trust calibration under real APC noise. This is the same qualitative regime targeted by AdaRL/FANsRL-style few-shot adaptation, while the direct AdaRL/FANsRL codebases are environment-specific RL systems rather than drop-in bus residual evaluators.

6.7 Auxiliary sampled rollout

The sampled BusSimEnv rollout is not primary evidence, but it checks whether the learned one-step action preferences are grossly inconsistent with executable line environments. Under a fair




Figure 7: Source-subset robustness. Removing data-source families reveals where the claim becomes weak.

seeded rollout over ten Austin/Halifax line episodes, weighted CFCMT reaches parity with the dense policy: mean reward is -1121.574 for weighted CFCMT and -1121.612 for $H2O^+$. The headway-error metric is slightly worse for weighted CFCMT, 588.2 seconds versus 585.6 seconds. Thus rollout supports a limited claim: after schema and reward alignment, CFCMT does not collapse in executable simulation, but the closed-loop advantage remains too small to serve as the main evidence.

7 Discussion

Why factored residuals help. The dense residual baseline has enough capacity to fit cross terms in source data, but the same flexibility can mix city-specific demand, speed, and headway effects. CFCMT restricts each output to a smaller mechanism-specific parent set. This constraint appears helpful in the strict cross-city setting: it reduces total MSE on all four targets and makes source weighting more effective.

What the source weighting does and does not show. Similarity weighting improves the mean strict leave-one-out ratio from 0.759 to 0.578 and wins over unweighted CFCMT on three of four targets. The Singapore target is the exception, where weighting is slightly worse than unweighted CFCMT. This suggests that static similarity is useful but imperfect; city descriptors built from open data may miss route-level operational differences that matter for transfer.

Policy evidence remains secondary. The one-step policy validation and sampled BusSimEnv rollout are useful checks, but they do not replace a closed-loop operational benchmark. The rollout result reaches reward parity with $H2O^+$ only after aligning BusSimEnv observations to the static feature schema and using identical line-level random seeds across policies. That correction is necessary for a fair comparison, but the resulting rollout margin is small. The strongest claim should therefore remain dynamics transfer plus policy-proxy improvement, not full policy dominance.

Static-derived validation is not operational validation. The central limitation is target construction. The experiments use full-route open static data and deterministic transition rules to create paired simulator/target transitions. This is valuable for reproducibility and reviewer scrutiny, but it is not the same as evaluating counterfactual holding actions on AVL/APC trajectories or a calibrated live SUMO rollout. The Austin raw APC/AVL slice partially addresses the trajectory-ground-truth concern for passive action-0 dynamics, but it still cannot identify how unobserved holding actions would have changed passenger and headway outcomes. The paper should therefore avoid claims such as “CFCMT is better on real bus systems”. A precise claim is: under open static-data-derived cross-city validation with little or no target-city calibration, and under passive real APC validation with few-shot adaptation, mechanism-factored residual transfer is more robust than dense residual correction.

How to strengthen the next version. The most valuable addition would be counterfactual or interventional validation on one city with vehicle-level AVL/APC trajectories, even if only for a subset of routes. The second priority is to expand executable line environments so sampled rollout can move from auxiliary sanity check to a stronger policy benchmark. The third priority is to add more cities, because city-level generalization is ultimately limited by the number of held-out cities.

8 Conclusion

We presented CFCMT, a causal-factored residual transfer approach for bus holding control under cross-city open-transit validation. Compared with a dense H2O⁺-style residual baseline, CFCMT improves strict leave-one-city-out dynamics on all four target cities, and source-city similarity weighting further improves the mean MSE ratio. One-step policy validation also favors CFCMT on most reward/regret and headway metrics and on all bunching-rate comparisons. At the same time, subset robustness shows that the advantage depends on source diversity, and generator robustness shows that not every perturbation preserves all-target wins. The Austin raw APC/AVL check adds a sharper boundary condition: source-only transfer is better than dense residual transfer but can still overshoot a strong state-informed passive baseline, while few-shot residual gating closes that gap. The resulting paper claim is intentionally bounded: CFCMT is a promising mechanism-factored alternative to dense residual correction for low-calibration cross-city transfer, but operational deployment claims require counterfactual control or field validation.

A Additional Tables

A.1 Target-construction audit

city	demand_evidence	traffic_evidence	lines	transitions	uncal_mse	line_mse_p25	line_mse_p50	line_mse_p75
Singapore	observed stop OD/PV monthly	observed speed bands	789	1243488	2256.5339	220.173	324.7508	541.4336
Austin / CapMetro	schedule proxy	schedule derived	234	310272	1301.3027	242.2482	615.8957	3043.8061
Halifax Transit	route APC apportioned	schedule derived	176	266448	1111.9223	199.2164	620.2908	1614.3016
MBTA	stop board/alight apportioned	schedule derived	940	969360	1216.3716	180.9625	459.6093	1040.0624

Table 5: Target-construction audit and route-level nontriviality check.

budget	splits	mean_no_cal_ratio	mean_cal_ratio	no_cal_wins	cal_wins	mean_calib_lines	oracle
0.0	4	0.5782	0.5782	4	4	0.0	False
0.01	4	0.5807	0.5801	3	3	3.5	False
0.05	4	0.5863	0.5831	3	3	23.0	False
0.1	4	0.6165	0.6094	3	3	50.5	False
0.25	4	0.6772	0.6596	3	3	132.75	False
1.0	4	0.7872	0.7242	3	3	534.75	True

Table 6: Target-route calibration-budget summary. Budget 1.0 is an in-sample oracle reference.

mechanism	targets	weighted_wins	mean_cfcmt_ratio	mean_weighted_ratio
demand_stop_reward	4	4	0.8197	0.7706
headway_gap	4	3	0.7586	0.5743
speed	4	4	0.8203	0.6684

Table 7: Per-mechanism strict leave-one-city-out error ratios.

A.2 Calibration budget

A.3 Per-mechanism errors

A.4 Mechanism and capacity ablations

family	wins_vs_h2oplus	mean_ratio_vs_h2oplus
cfcmt_shared_sparse	4	0.7502
cfcmt_no_action_interaction	4	0.7583
cfcmt_full	4	0.7587
dense_matched_sparse	3	0.7716
dense_all_parents_no_pairwise	3	0.9332
cfcmt_action_time_only	1	7.3401
cfcmt_random_mechanism_grouping	1	7.5744

Table 8: Mechanism and capacity ablations. Lower mean ratios indicate better transfer relative to H2O⁺.

target_city	lines	cfcmt_ratio	cfcmt_ratio.ci95	weighted_ratio	weighted_ratio.ci95
Singapore	789	0.9964	[0.991, 1.000]	0.9986	[0.993, 1.002]
Austin / CapMetro	234	0.7375	[0.670, 0.817]	0.3985	[0.354, 0.447]
Halifax Transit	176	0.7166	[0.529, 0.871]	0.5638	[0.310, 0.783]
MBTA	940	0.5844	[0.520, 0.643]	0.352	[0.257, 0.438]

Table 9: Route-level bootstrap confidence intervals under strict leave-one-city-out dynamics validation.

dataset	transitions	routes	stops	trips	date_min	date_max	h2oplus_vs_passive	weighted_vs_passive	cfcmt_ratio	weighted_ratio
CapMetro APC Raw July 2021 - December 2021	5771472	68	2315	421255	2021-07-01	2021-12-30	10.1614	3.352	0.5281	0.3299
CapMetro APC Raw August 2016	5796060	72	2576	220011	2016-08-21	2017-01-07	10.3665	3.3452	0.5173	0.3227

Table 10: External Austin raw APC/AVL validation. The observed data provide passive action-0 stop-event transitions, not counterfactual holding actions.

A.5 Route-bootstrap confidence intervals

A.6 External real APC/AVL validation

A.7 Configured six-split dynamics validation

A.8 Source-weighting sensitivity

A.9 Source-size sensitivity and single-source transfer

A.10 Generator robustness

A.11 Source-subset robustness

A.12 Auxiliary sampled rollout

References

- Francesco Alesiani and Konstantinos Gkiotsalitis. Reinforcement learning-based bus holding for high-frequency services. In *2018 21st International Conference on Intelligent Transportation Systems*, pages 3162–3168, 2018. doi: 10.1109/ITSC.2018.8569473.
- Yoshua Bengio, Tristan Deleu, Nasim Rahaman, Nan Rosemary Ke, Sébastien Lachapelle, Olexa Bilaniuk, Anirudh Goyal, and Christopher Pal. A meta-transfer objective for learning to disentangle causal mechanisms. In *International Conference on Learning Representations*, 2020.
- Carlos F. Daganzo. A headway-based approach to eliminate bus bunching: Systematic analysis and comparisons. *Transportation Research Part B: Methodological*, 43(10):913–921, 2009. doi: 10.1016/j.trb.2009.04.002.
- Benjamin Eysenbach, Shreyas Asawa, Shreyas Chaudhari, Sergey Levine, and Ruslan Salakhutdinov. Off-dynamics reinforcement learning: Training for transfer with domain classifiers. In *International Conference on Learning Representations*, 2021.
- Fan Feng, Biwei Huang, Kun Zhang, and Sara Magliacane. Factored adaptation for non-stationary reinforcement learning. *arXiv preprint arXiv:2203.16582*, 2022.

dataset	budget	cal.lines	cal.transitions	source_only_vs_passive	source_plus_target_vs_passive	target_only_vs_passive	residual_gate_vs_passive
CapMetro APC Raw July 2021 - December 2021	0.0	0	0	3.352	3.352	-	3.352
CapMetro APC Raw July 2021 - December 2021	0.01	1	745	3.352	0.9803	1.0662	0.97
CapMetro APC Raw July 2021 - December 2021	0.05	8	293425	3.3466	0.93	0.8982	0.9701
CapMetro APC Raw July 2021 - December 2021	0.1	12	650668	3.3437	0.9207	0.9087	0.9699
CapMetro APC Raw July 2021 - December 2021	0.25	24	1537062	3.3381	0.911	0.8933	0.9699
CapMetro APC Raw August 2016	0.0	0	0	3.3452	3.3452	-	3.3452
CapMetro APC Raw August 2016	0.01	1	5088	3.3452	0.9515	0.9765	0.9567
CapMetro APC Raw August 2016	0.05	8	512496	3.3556	0.9028	0.8749	0.9551
CapMetro APC Raw August 2016	0.1	14	922153	3.3472	0.8963	0.8784	0.9549
CapMetro APC Raw August 2016	0.25	26	1620309	3.3277	0.8903	0.879	0.9556

Table 11: Few-shot Austin raw APC/AVL adaptation. Ratios are total MSE relative to the state-informed passive no-correction baseline on held-out Austin line groups.

split	target_city	train_transitions	target_transitions	h2oplus_total_mse	cfcm_total_mse	cfcm_vs_h2oplus_ratio	weighted_cfcm_vs_h2oplus_ratio	weighted_cfcm_vs_unweighted_ratio
source.open.na.to.singapore	Singapore	1546080	1243488	2254.8548	2246.6693	0.9964	0.9986	1.0022
source.singapore.austin.halifax.to.mbt	MBTA	1820208	969360	318.9925	186.4177	0.5844	0.352	0.6022
leave_one_city_out.all:singapore.it.all	Singapore	1546080	1243488	2254.8548	2246.6693	0.9964	0.9986	1.0022
leave_one_city_out.all:austin.capmetro.all	Austin / CapMetro	2479296	310272	91.9724	67.8254	0.7375	0.3985	0.5404
leave_one_city_out.all:halifax.transit.all	Halifax Transit	2523120	266448	97.1175	69.5992	0.7166	0.5638	0.7867
leave_one_city_out.all:mbta.all	MBTA	1820208	969360	318.9925	186.4177	0.5844	0.352	0.6022

Table 12: Configured cross-city dynamics splits. This table includes repeated target cities and is treated as appendix evidence rather than the primary independent-sample result.

Biwei Huang, Fan Feng, Chaochao Lu, Sara Magliacane, and Kun Zhang. AdaRL: What, where, and how to adapt in transfer reinforcement learning. In *International Conference on Learning Representations*, 2022.

JB Lanier, Kyungmin Kim, Armin Karamzade, Yifei Liu, Ankita Sinha, Kat He, Davide Corsi, and Roy Fox. Adapting world models with latent-state dynamics residuals. In *International Conference on Machine Learning*, 2025.

Dong Liu, Feng Xiao, Jian Luo, and Fan Yang. Deep reinforcement learning-based holding control for bus bunching under stochastic travel time and demand. *Sustainability*, 15(14):10947, 2023. doi: 10.3390/su151410947.

Pablo Alvarez Lopez, Michael Behrisch, Laura Bieker-Walz, Jakob Erdmann, Yun-Pang Flötteröd, Robert Hilbrich, Leonhard Lücken, Johannes Rummel, Peter Wagner, and Evamarie Wießner. Microscopic traffic simulation using SUMO. In *2018 21st International Conference on Intelligent Transportation Systems*, pages 2575–2582, 2018. doi: 10.1109/ITSC.2018.8569938.

Mitsuhiko Nakamoto, Yuexiang Zhai, Anikait Singh, Max Sobol Mark, Yi Ma, Chelsea Finn, Aviral Kumar, and Sergey Levine. Cal-QL: Calibrated offline RL pre-training for efficient online fine-tuning. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

Haoyi Niu, Shubham Sharma, Yiwen Qiu, Ming Li, Guyue Zhou, Jianming Hu, and Xianyuan Zhan. When to trust your simulator: Dynamics-aware hybrid offline-and-online reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 35, pages 36599–36612, 2022.

Haoyi Niu, Tianying Ji, Bingqi Liu, Haocheng Zhao, Xiangyu Zhu, Jianying Zheng, Pengfei Huang, Guyue Zhou, Jianming Hu, and Xianyuan Zhan. H2O+: An improved framework for hybrid offline-and-online reinforcement learning with dynamics gaps. *arXiv preprint arXiv:2309.12716*, 2025.

Jiawei Wang and Lijun Sun. Dynamic holding control to avoid bus bunching: A multi-agent deep reinforcement learning framework. *Transportation Research Part C: Emerging Technologies*, 116: 102661, 2020. doi: 10.1016/j.trc.2020.102661.

temperature	floor	wins_vs_h2oplus	wins_vs_unweighted_cfcmt	mean_ratio_vs_h2oplus	mean_ratio_vs_unweighted_cfcmt
0.5	0.0	6	4	0.6025	0.7449
0.5	0.05	6	4	0.6033	0.746
0.5	0.1	6	4	0.6041	0.747
1.0	0.0	6	4	0.6101	0.7554
1.0	0.05	6	4	0.6106	0.756
1.0	0.1	6	4	0.611	0.7566
2.0	0.0	6	4	0.6145	0.7614
2.0	0.05	6	4	0.6147	0.7617
2.0	0.1	6	4	0.6149	0.762

Table 13: Source-weighting sensitivity over temperature and floor.

source_size	splits	cfcmt_wins_vs_h2oplus	mean_ratio_vs_h2oplus
1	12	6	0.9272
2	12	8	0.8705
3	4	4	0.7587

Table 14: Mean cross-city ratio by number of source cities.

Jiawei Wang and Lijun Sun. Multi-objective multi-agent deep reinforcement learning to reduce bus bunching for multilane services with a shared corridor. *Transportation Research Part C: Emerging Technologies*, 156:104309, 2023. doi: 10.1016/j.trc.2023.104309.

Yiguang Xuan, Juan Argote, and Carlos F. Daganzo. Dynamic bus holding strategies for schedule reliability: Optimal linear control and performance analysis. *Transportation Research Part B: Methodological*, 45(10):1831–1845, 2011. doi: 10.1016/j.trb.2011.07.009.




Figure 8: Single-source transfer matrix. Cells report CFCMT/H2O total-MSE ratio.

scenario	splits	weighted_wins_vs_h2oplus	unweighted_wins_vs_h2oplus	weighted_wins_vs_unweighted	mean_unweighted_ratio_vs_h2oplus	mean_weighted_ratio_vs_h2oplus	mean_weighted_ratio_vs_unweighted
baseline	4	4	4	3	0.7587	0.5782	0.7329
headway_gap_bias_125	4	3	3	3	0.7587	0.5768	0.7301
demand_stop_reward_bias_125	4	4	4	3	0.7587	0.5802	0.737
speed_bias_150	4	4	4	3	0.7589	0.5786	0.7332
mixed_mechanism_bias	4	3	3	3	0.7588	0.5767	0.7295
measurement_noise_10pct	4	4	4	3	0.778	0.614	0.7648
measurement_noise_25pct	4	4	4	3	0.8391	0.7255	0.851

Table 15: Generator-bias and measurement-noise robustness.

subset	cities	splits	weighted_wins_vs_h2oplus	unweighted_wins_vs_h2oplus	mean_unweighted_ratio_vs_h2oplus	mean_weighted_ratio_vs_h2oplus	mean_weighted_ratio_vs_unweighted
all_four_attrict_leave_one_out	Singapore, Austin / CapMetro, Halifax Transit, MBTA	4	4	4	0.7587	0.5782	0.7329
exclude_singapore_attrict_traffic_only	Austin / CapMetro, Halifax Transit, MBTA	3	1	1	1.0025	0.9984	0.9964
exclude_austin_attrict_demand_proxy	Singapore, Halifax Transit, MBTA	3	3	3	0.7898	0.6626	0.8124
north_american_observed_ridership_pair	Halifax Transit, MBTA	2	0	0	1.0421	1.0421	1.0

Table 16: Source-subset robustness.

policy	episodes	mean_reward	mean_headway_abs_error	mean_hold_seconds
no_hold	10	-1138.411	601.9113	0.0
fixed_30	10	-1151.2901	598.4459	30.0
h2oplus_dense_policy	10	-1121.612	585.5526	3.117
cfcmt_mechanism_policy	10	-1122.2105	589.1995	3.4897
cfcmt_similarity_weighted_policy	10	-1121.574	588.2357	3.2859

Table 17: Auxiliary sampled BusSimEnv rollout. This result is not primary evidence because only a small number of runnable line environments are available.